

Akshay Dhar

Bengaluru, India

akshaykumardhar14@gmail.com

akshaykumardhar.github.io, linkedin.com/in/akshaydh, github.com/AKSHAYKUMARDHAR

AI Product Manager who takes AI from problem to production and gates every release on an evaluation. At Catalysk I owned a transaction categorization engine that auto-categorizes when confident and routes uncertain rows to human-in-the-loop review, cutting uncategorized transactions from 39% to 13%. At 9fin I shipped the Gemini API into production automation with self-written prompts, and owned an ML crawler delivering £390K in annual savings. I prototype and build with Claude Code and Cursor.

Core Competencies

- **AI Product Ownership:** problem diagnosis, solution design, prototyping, shipping and iterating an AI categorization engine in production; quality and precision trade-off calls backed by evidence.
- **Evaluations & Release Bars:** golden sets, held-out sets, benchmarks and confidence gating; embedding model bake-offs (MiniLM, E5, BGE).
- **LLMs, AI Agents & RAG:** Gemini API in production with self-written prompts; LLM agent with MCP tool calling, RAG on pgvector and prompt-injection guardrails (personal project).
- **Building:** Claude Code and Cursor; Python (Pandas), SQL/PostgreSQL, FastAPI, n8n, REST APIs, config-driven YAML pipelines, 90+ automated data quality checks.
- **Domains:** fintech (UPI/VPA, P2P, institutional credit), ESG and sustainability, energy, edtech.

Professional Experience

- **Catalysk, AI Sustainability Analytics** **Bengaluru, India**
Data Consultant *Apr 2026 - Sep 2026*
 - Architected and led the build of a multi-stage categorization engine with a small team (rules, merchant extraction, an SBERT fallback and the evaluation harness), owning it from design through production.
 - Cut uncategorized transactions from 39% to 13% by redesigning UPI (VPA) and P2P categorization logic feeding client ESG reporting.
 - Gated auto-categorization at ~84% model confidence, routing uncertain rows to analyst review (human-in-the-loop) to prioritise precision over vanity coverage.
 - Defined the release bar with a 388-row UAT golden set and a 1,400-row benchmark across 7 bank statements, gating every model and taxonomy change.
 - Lifted category accuracy 5 points with an embedding model bake-off (MiniLM, E5, BGE) against the golden set, rejecting a slower model that only looked better on paper.
 - Grew model coverage from 21 to 36 categories and 10K to 17K training examples by retraining SBERT and merging a second tenant's taxonomy, with zero golden-set regression.
 - Cut new banking partner onboarding from days to hours with a config-driven YAML framework and one ingestion contract for CSV, Excel, JSON and REST, so analysts change rules without engineering.
 - Recovered a catch-all bucket at ~30% of volume by diagnosing a silent regression that disabled SBERT and 4 downstream stages; cut processing from hours to minutes per 1,000 rows with embedding caching.
 - Shipped 90+ automated data quality checks across 5 financial schemas, and an 18-table PostgreSQL observability layer with live Metabase dashboards replacing manual Excel review.

- 9fin, Institutional Financial Technology Platform**
Senior Data Associate

Belfast, UK
Jan 2025 - Sep 2025

 - Delivered £390K annual cost savings and 32% faster client delivery by owning a production ML market-news crawler, resolving incidents and adding verification controls.
 - Accelerated client onboarding by 15% by shipping the Gemini API into production automation for post-acquisition schema alignment, with self-written prompts.
 - Mentored 15 analysts on data quality standards, verification workflows and tooling.
- 9fin, Institutional Financial Technology Platform**
Data Associate

Belfast, UK
May 2024 - Dec 2024

 - Maintained 100% data accuracy against delivery SLAs across 140+ institutional credit instruments.
 - Built real-time KPI and incident dashboards (Power BI, Notion) for product, engineering and non-technical stakeholders.
- JamBird Marketing**
Data Science Intern

Belfast, UK
Jun 2023 - Sep 2023

 - Deployed an NLP text classification model at 87% accuracy that automated insight generation and replaced a manual review process.
 - Engineered features across 3M+ data points and presented findings to senior stakeholders.
- BYJU'S**
Data Analyst

India
Aug 2020 - Jul 2021

 - Implemented ML learning-path recommendations, improving student segmentation accuracy by 15%.
 - Optimised PySpark ETL pipelines, improving data processing time by 20%.

Selected Projects

- Personal project, github.com/AKSHAYKUMARDHAR/UPI-Triage-Agent**
UPI Transaction Triage Agent

Oct 2026

Built with Claude Code: an LLM agent with MCP tool calling, RAG on pgvector, a human review queue, n8n and FastAPI, and a prompt-injection guard. Agent + RAG reached 99.6% automation at 99.6% precision on a 250-row golden set, and 98.0% precision on 57 never-seen rows with 0% of ambiguous rows auto-posted wrong (44.4% without RAG).
- Personal project**
End-to-End Stock Price Forecasting Pipeline

Feb 2026 - Mar 2026

Live market data via API into PostgreSQL, with ARIMA and LSTM multi-step forecasting.
- Energia Group, MSc industry project**
Electricity Market Price Forecasting

2023

3-model forecasting system (VAR, STL, Random Forest) for trading decisions, surfaced through executive Tableau dashboards.

Education

- Queen's University Belfast**
MSc Data Analytics

Belfast, UK
2022-2023
- Assam University Silchar**
BSc Computer Science

Silchar, India
2017-2020

Certifications, Languages & Awards

- Oracle Cloud Infrastructure AI Foundations Associate; Google Data Analytics Professional Certificate; BCG GenAI Job Simulation and BCG Data Science Job Simulation (Forage).
- English (full professional), Hindi (native), Assamese (native), Bengali (native).
- Best Project, Irish Intervarsity (2023).

References

- Available upon request.