

Is This a Scam?

A scam checker for India in Hindi, Bengali and English that says "can't tell" instead of guessing.

Akshay Dhar · akshaykumardhar14@gmail.com · is-this-a-scam.onrender.com · github.com/AKSHAYKUMARDHAR/Is-This-A-Scam

MY ROLE	WHEN	STACK
Solo: research, PRD, evaluation design and the build, with Claude Code.	Oct 2026	Python, FastAPI, Gemini (free tier) or Claude, vanilla JS, Docker, Render
1 / 114 false alarms on never-seen genuine messages, over two release runs	0 / 96 never-seen scams wrongly cleared	
0 / 50 prompt-injection attacks that changed a verdict	Blocked twice by my own launch gate: ambiguous messages, then 1 false alarm	

The problem

Indians lost **₹19,813 crore to cyber fraud in 2025**, and 77% of it went to fake investment schemes (I4C data). Telecom and inbox filters act silently, and only on what they are sure about. What gets through is the ambiguous message or call a person is staring at, wondering "is this real?" Nobody gives that person a clear, explained answer in their own language.

The market is crowded: Airtel and Google screen billions of messages, and Truecaller and WhatsApp are adding scam checks. So the product can't win on detection volume. It wins on three things no checker I reviewed combines: **honest uncertainty**, **coverage of investment pitches and phone calls** (the scams behind 85% of losses, which link filters rarely see), and **Bengali** alongside Hindi and English.

Who it's for

- **The unsure recipient**, often 45 to 70, who gets a "your KYC expires today" SMS or a caller claiming to be police.
- **The family guardian** who handles a parent's phone safety and shares answers into family WhatsApp groups. They are the growth channel.
- **The targeted investor** added to a "stock tips" group, the 77%-of-losses case.
- **The small shopkeeper** shown a "payment successful" screenshot, where the right answer is "a screenshot can never prove payment".

What I built

A mobile-first web app with no login. Paste a message, upload a screenshot, or answer five questions about a call, and get one of three verdicts: **Likely scam**, **Can't tell, verify it yourself**, or **No scam signs found**. It never says "safe".

Input

message, screenshot
or call answers

Rules + injection guard

11 hard rules decide
what is always fraud

Masking

card, account, ID,
phone and OTP
numbers hidden

2 model runs

must agree, or the
answer is "can't tell"

Verdict card

red flags quoted from the message, fixed next steps

क्या यह स्कैम है?

EN हिंदी বাংলা

कोई संदिग्ध संदेश पेट्ट करे, स्क्रीनशॉट डालें या कॉल के बारे में बताएं। अपनी भाषा में साफ़ जवाब पाएं।

मैंने ऐसे भेज दिए या OTP बता दिया

संभवतः स्कैम

प्रकार: नकली बैंक, KYC या कार्ड संदेश

यह संदेश बैंक बनकर आपका OTP मांग रहा है। असली बैंक कभी OTP नहीं मांगते।

हमें मिले खतरे के संकेत

"आया हुआ OTP इस नंबर पर भेजें"

OTP, PIN, CVV या पासवर्ड बताने को कहा गया है। बैंक और असली कंपनियां ये कभी नहीं मांगतीं।

"sbi-kyc-update.xyz"

लिंक छोटा या अजीब देब पता है, जिससे पता नहीं चलता कि वह कहाँ ले जाएगा।

"आज रात खाता बंद"

जल्दबाज़ी का दबाव डालकर सोचने का समय नहीं देता।

अब क्या करें

1. ऐसे न भेजें, किसी लिंक पर क्लिक न करें, और OTP, UPI पिन या कार्ड की जानकारी किसी को न दें।

2. भेजने वाले को ब्लॉक करें और sancharsaathi.gov.in पर चक्षु में इस संदेश या कॉल की शिकायत करें।

3. अगर आपने ऐसे भेज दिए हैं या जानकारी दे दी है, तो तुरंत 1930 पर कॉल करें या cybercrime.gov.in पर शिकायत करें।

सुरक्षित तरीके से कैसे जाँचें

अपने बैंक की आधिकारिक ऐप खोलें या कार्ड के पीछे लिखे नंबर पर कॉल करें। बैंक KYC के लिए लिंक नहीं भेजते और OTP नहीं मांगते।

यह कोई गारंटी नहीं है। मदद के लिए 1930 पर कॉल करें।

WhatsApp पर शेयर करें

क्या यह मददगार था?

हाँ

नहीं

क्या इससे आप ऐसे भेजने या लिंक खोलने से रुके?

Is This a Scam?

EN हिंदी বাংলা

Paste a suspicious message, upload a screenshot, or describe a call. Get a clear answer in your language.

I already paid or shared my OTP

This could be a real bank alert, but the short link hides where it goes. Check in your bank's own app instead.

Red flags we found

"bit.ly/3stmnt-sep"

The link uses a short or unusual web address that hides where it goes.

What looks normal

No request for an OTP, PIN or payment.

Can't tell. Verify it yourself

What to do now

1. Don't act on this message yet: don't pay, click or reply.

2. Check directly with the organisation through its official app, or a number you look up yourself. Never use a number or link from the message.

3. Never share an OTP, UPI PIN or CVV, whoever asks.

How to check safely

Open your bank's official app, or call the number on the back of your card. Banks never send links for KYC or ask for OTPs.

This is not a guarantee. For help, call 1930.

Share on WhatsApp

Was this helpful?

Yes

No

Did this stop you from paying or clicking?

Yes

This verdict is wrong

এটা কি স্ক্যাম?

EN हिंदी বাংলা

সন্দেহজনক বার্তা পেস্ট করুন, স্ক্রিনশট দিন বা কলের কথা লিখুন। নিজের ভাষায় পরিষ্কার উত্তর পান।

আমি টাকা পাঠিয়েছি বা OTP বলে দিয়েছি

বার্তা বা স্ক্রিনশট

ফোন কল

তারা নিজেকে কে বলে পরিচয় দিল?

পুলিশ, সিবিআই, কাষ্টমস বা আদালত

ব্যাংক বা RBI

TRAI বা ফোন কোম্পানি

কুরিয়ার কোম্পানি

পরিবারের কেউ বা বন্ধু

কোম্পানি বা HR

অন্য কেউ / জানি না

তারা আপনাকে কী করতে বলল? যতগুলো প্রযোজ্য বেছে নিন।

টাকা পাঠানো

OTP, পিন বা কার্ডের তথ্য জানানো

কোনো অ্যাপ ইনস্টল করা

আধার বা ব্যাংকের তথ্য দেওয়া

কলে থাকা বা ক্যামেরা চালু রাখা

এখনও কিছু না

তারা কি হুমকি দিয়েছে?

শ্রেয়তর বা পুলিশ কেস

অ্যাকাউন্ট, কার্ড বা সিম বন্ধ করা

পরিবারের কারও ক্ষতি

কোনো হুমকি নেই

এটা কি ভিডিও কল ছিল, বা তারা গোপন রাখতে বলেছিল?

হ্যাঁ

না

তারা কি টাকা কোনো 'নিরাপদ' বা 'RBI' অ্যাকাউন্টে সরাসরি বলেছিল?

হ্যাঁ

না

তারা আর কী বলেছিল? (ঐচ্ছিক)

নিজের ভাষায়...

এখনই যাচাই করুন

Also in the product: an "I already paid" button that goes straight to the 1930 helpline steps, a WhatsApp share card that never includes the original message, a verdict-first vs explanation-first A/B split, feedback buttons, and an event log that stores no message text. When the free model quota runs out, messages still get the fixed rules' answer instead of an error.

Product decisions

- **Three verdicts, never "safe".** One wrong "scam" on a genuine bank message costs more trust than ten honest "can't tell" answers.
- **Rules decide the cases that are always fraud.** Banks never ask for an OTP, and you never need a PIN to receive money. Every hard rule is tested against genuine look-alikes: OTP messages that say "do not share", delivery OTPs you *do* give the agent, police advisories about digital arrest.
- **Self-consistency instead of self-reported confidence.** Two runs must agree. The PRD planned three; the golden set showed the third never changed a verdict (128 of 128), so I cut it and saved a third of the cost.
- **No ask, no "Likely scam" from the model alone.** A first-contact opener ("this is my new number, save it") asks for nothing yet, so nothing can be lost yet. The card says "can't tell" and what to watch for in the next message. I added this rule after my first release run failed on exactly these messages.
- **The injection guard sits outside the model,** a lesson from my UPI agent, where a system prompt alone did not stop an injection.
- **Advice is never generated.** Next steps, numbers and links are fixed templates in three languages; the model's quotes are checked against the message and dropped if they don't appear.

How I tested it

I froze a 200-message held-out set in git **before writing any rule or prompt**, tuned on a separate 200-message golden set, and ran the held-out set once on the frozen configuration. When it failed, I fixed the problem on the golden set only, wrote **50 fresh messages and froze them before the fix existed**, and ran those once. The messages are in Hindi, Bengali and English, including Hindi and Bengali typed in Latin script, plus phone-call descriptions.

Release run	False alarms	Scams cleared	Scams caught	Ambiguous → can't tell
Run 1 · 200 messages	0 / 100	0 / 80	100%	50%
Run 2 · 50 fresh messages, after the fix	1 / 14	0 / 16	87.5%	85%

In run 1, the rules alone caught 47.5% of scams with no false alarm, and requiring two model runs to agree removed the one false alarm a single run made.

Release gate, run 2

- ✓ Ambiguous messages get "can't tell" $\geq 70\%$: 85% (run 1: 50%)
- ✓ Scams cleared $\leq 5\%$: 0 of 16
- ✓ Scam type correct $\geq 85\%$: 100%
- ✓ Prompt-injection suite, zero verdict flips: 0 of 50
- ✗ False alarms $\leq 2\%$ of genuine messages: 1 of 14
- ... Native-speaker review of the explanations: pending

What failed, and what's next

My own gate blocked the launch twice. In run 1 the model was too sure on first-contact messages with no request yet ("Maa, this is my new number, save it"), so I added a rule: the model alone can't call a message a scam unless it asks for something risky (pay, open a link, call a number, share a code, install an app). Run 2 met the ambiguous target but called one genuine Hindi SBI debit alert a scam ("if this wasn't you, call 1800 1234"), a format fraudsters copy. So the live site stays a demo, not a launched product.

Two honest caveats from run 2. On those fresh messages the model alone would already have reached 80% on ambiguous messages, so the run shows the target is met, not that my rule is what meets it; the rule also turned two "press 9"-style scams into "can't tell", though neither was cleared. And with only 14 genuine messages, one mistake is 7%: my second test set was too small to test a 2% bar fairly.

Next: treat bank alerts that point to an official helpline as a known genuine pattern, re-test on at least 100 genuine messages, preferably real ones, and run the second-labeller and native-speaker reviews, whose kits are in the repo.

What I learned

- A held-out set frozen before the rules is what makes the numbers worth trusting, and every fix needs a fresh one.
- Check the counterfactual. Re-scoring the same cached answers without my fix showed how little it did on the second set.
- Size a test set for the bar it tests: 14 genuine messages can't measure a 2% false-alarm rate.
- Measure the expensive part. The third model run looked prudent and did nothing.
- Free tiers shape evaluation plans: 500 requests a day per project decided how much I could test, and in what order.

All 500 test messages are hand-written, modelled on publicly reported scam patterns and real message formats, by the same person who wrote the rules; results on real messages may differ. Every number comes from a logged run in the repo.